

The Nauti Yacht — Owner Console Manual

For anyone using the admin dashboard at **thenautiyacht.com/admin**. Ask the owner for the passcode — it is not in this manual.

Last updated 18 September 2026 (fourth revision).

How the console is laid out

Six groups run along the top. Clicking a group reveals its tabs underneath.

Group	Tabs
Overview	The day at a glance, the bank balance, the board, and the crew
Bookings	Contacts · Bookings · Availability
Money	Income & expenses · Reconciliation · Tax Report · Gift certificates · Subscriptions & bills
Marketing	Media · Media Drafts · Comments · Messages · Testimonials · Photo Requests
Setup	Packages & pricing · Add-ons · Coupons
Boat	Maintenance

A number in brackets after a tab name means **something is waiting on you**. No number means nothing needs doing — it does not mean the tab is empty. Testimonials holds every approved review and still shows no number, because none of them need a decision.

Bookings

Contacts

People, and how to reach them. This tab was called Inquiries and used to list them; it no longer does, because **every reservation lives in Bookings** — inquiry, booked or completed, whether it came from Boatsetter, GetMyBoat, the website or a text message. A lead that showed in one place and not the other is exactly how one went missing.

Three lists, all collapsed until you open them:

- **Everyone we can contact** — people, not trips, so a repeat guest appears once. The count is the ones with a phone or an email; anyone with neither cannot be asked for a review or told about a glow night.
- **Extra guest contacts** — people who were on somebody else's charter and whose number is worth keeping. Deliberately not counted as inquiries.

- **Crew list** — emails captured from /glow and the on-boat QR code. This is the list to mail when a date is set. **Copy mailable emails** sits on the header so you do not have to open the panel to use it. To honour an unsubscribe set that person's status to lapsed and they drop out of the copy.

The two lower panels are subsets of the top one, not additions to it. Their headings say so — "*2 of the 30 above*" — because the obvious arithmetic is wrong: the contacts list already folds in crew-list signups and extra contacts, and dedupes a person to one row however many lists they appear on. Adding 30 and 2 and 2 counts four people twice.

Bookings

Every charter, from any source — Boatsetter, GetMyBoat, the website, cash, Zelle. Add one with the form above the table.

Three questions, three fields

These used to be one and a half fields, and the books paid for it. Each booking now answers them separately, and they are genuinely different questions:

field	the question	options
Lead source	Where did the enquiry come from?	Website · Text / WhatsApp · Phone · Walk-up · Instagram · Facebook · Repeat guest · Referral from a friend · Boatsetter · GetMyBoat · AI search (ChatGPT) · Other
Booking channel	Who took and processed it?	Boatsetter · GetMyBoat · Website · Direct
How paid	How did the money actually arrive?	Unpaid · Stripe (card) · Cash · Cash App · Zelle · Venmo · PayPal · Boatsetter payout · GetMyBoat payout · Gift certificate

A GetMyBoat booking can come from a guest who found us on Instagram. That is the interesting fact, and until now there was nowhere to put it.

"How paid" is the one the system cannot work out for itself. Stripe paying is the single exception — that gets set on its own. Everything else has to be said by you, because nothing anywhere proves cash changed hands except you saying so. Leave it blank and the income row says so rather than guessing.

That guess is why this exists. The booking channel used to be read as the payment method — anything taken directly was assumed to be cash — so the first card payment on a text booking would have been filed in the cash column, the hardest one to reconcile and the easiest place to lose a number.

The old `other` channel is now **Direct**, which is what those bookings always were: ones we took ourselves, by text, on the phone or at the dock.

Marking a booking "completed" now writes its income row automatically. That was the single biggest hole in the system: the two records were only ever joined by hand, and six charters' income went missing that way. Two rules govern it:

- It will not create a second row if one already exists, so re-saving a completed booking is safe.

- **No price means no income row.** It refuses to guess, because a made-up number in the ledger is worse than an obviously missing one.

A booking paid through the website checkout now appears here on its own, created from the Stripe webhook as **booked** (not completed — the trip has not happened yet).

One charter, two rows — and they now stay in step

A website checkout leaves **two** records for the same charter: the Inquiry the guest filled in, and the mirror booking above that blocks the date on the calendar. Nothing used to join their statuses.

So marking a charter **completed** used to leave its inquiry record still reading **booked and paid** — which is what happened to Oscar RoblesGil's 6 September charter. He was the first website checkout to produce such a pair, which is the only reason nobody had seen it; every one after him would have done the same.

Changing the status updates both records, in both directions. Since 11 Sep 2026 there is only one place to do it — the Bookings list — but the rule still matters, because the two rows are what the calendar and the ledger each read. Platform bookings — Boatsetter, GetMyBoat, cash — have no inquiry behind them and are unaffected.

Either record completing writes the income row, by the same rule as above. That was worth being careful about: without it, which screen you happened to use would have decided whether a charter's money got recorded.

The same inquiry sent twice

Sarah Griffith sent the same glow-night inquiry **three minutes apart** on 12 September. Nothing was wrong with the first one. The form simply reset itself and flashed a message for two and a half seconds, so thirty seconds later her screen looked exactly like a form she had never filled in — and she filled it in again. You got two rows and she got two acknowledgement emails for one charter.

Two separate things now stop that.

The website confirms it on screen and leaves it there. The form is replaced by a panel that reads back what she asked for — charter, boat, date, guests, quoted price — says where her confirmation email is going, and says plainly that there is nothing else to do and no need to send it again. It does not time out. There is a *Send another inquiry* button for anyone who genuinely wants a second.

And the site refuses to write the duplicate anyway. The panel only helps somebody whose browser still has the page; a second tab, a phone after a laptop, or a form reloaded later never sees it. So a submission that matches an inquiry already sitting in the console is folded into that one instead of making a new row.

Counts as the same inquiry	same email address, same requested date, same package, same boat
Within	30 minutes of the first one
Only if	you have not touched the first one yet — it is still new

Party size, phone number and message are deliberately **not** part of that test, because those are exactly what somebody corrects on a second try, and a correction is still the same charter.

The boat is part of it, though, and that is the interesting one. A group too big for one deck books *two* — same person, same date, same package, two vessels, very plausibly half an hour apart while they count heads. Leaving the boat out of the test meant the second inquiry was folded into the first and the different vessel filed as a "corrected boat": not a duplicate prevented, half a booking gone. Including it costs the opposite case — somebody who changes their mind about the boat and resubmits gets a second row, which is the old behaviour and one click to delete. Those two are not comparable, so the tie goes to never losing the booking.

The 30 minutes is not a guess at how fast people double-click — it is how long "I am trying to book this one charter" lasts as a single sitting. Someone who fills the form, goes to check a date with a friend and comes back twenty minutes later is on the same errand. It is deliberately not a whole day: two inquiries for the same date a week apart are worth seeing separately, because by then something has changed.

What you get told. Nothing, if they changed nothing — a double-click is not news, and the email you already have is still correct. If they *did* change something, the row is updated to the newer details and you get one email whose subject begins **"Updated inquiry"** and whose first line says *NOT A NEW INQUIRY* and names what changed. The guest gets no second acknowledgement either way: a second email about one charter is what causes the "do I have two bookings?" phone call this exists to prevent.

Three things it will not do. It will not merge two inquiries from the same person for *different* dates, *different* packages or *different* boats. It will not touch an inquiry you have already booked, lapsed or cancelled — if somebody asks again after a cancellation, that is new intent and gets its own row. And an inquiry with **no date** is never deduplicated, because there is nothing to identify it by, and wrongly merging two "someday" inquiries loses a real lead — which is much worse than a duplicate you delete.

Where they came from, and how they paid

These are two different questions and the books used to answer them in one column. A guest can find you on Instagram and hand you cash. They can come through the website and pay by Venmo. Boatsetter and GetMyBoat always pay the same way — a payout straight to the bank — but everything else is free to combine however it likes.

	Says	Lives on
Lead source	where the guest came from	the booking
How paid	how the money arrived	the booking
Ledger origin	how the money arrived	each income or expense row

The ledger's origin had been doing both jobs, so income was filed under **Instagram** (\$775), **Friends** (\$928.06), **Yolo Lake Conroe** (\$2,400) and **Website** — none of which is a way of paying. Worse, the income dropdown offered only sources, so there was *literally no way* to record that a charter was paid in cash. Both dropdowns are now built from the payment-method list, so one cannot gain an option the other lacks.

Corrected 13 September 2026, across 24 rows. Totals did not move — this changed what the money is filed under, never how much of it there is.

Cash App is not Cash. It leaves a statement that can be reconciled and can charge a fee; cash in a hand does neither, and the whole point of the field is telling those apart. It files to the ledger as *Cash App Statement*, the spelling your expense rows already used.

One charter can be paid two ways. Rheya Palmer bought two hours by card and topped up to three by Venmo; Nagdy's extra hour came by Zelle on top of a Boatsetter payout. The booking's *How paid* holds the principal method and each actual payment is its own ledger row, so the books can still say how every dollar arrived.

One thing spelled two ways is two things. The expense side carried "Gmail" and "Gmail Statement" for the same email receipts, and "T-Mobile" and "Tmobile Statement" for the same phone bill — four totals split in half, with neither figure the real one. Merged 13 September 2026; the totals did not move, because renaming money never should.

Two origins were carrying a *detail* rather than a method — which bank a Venmo came out of, who the cash went to. Those moved into the note, where they are still readable, instead of splitting a total.

The morning check now watches for both faults: an origin the console cannot display (which would silently refile a row the next time you edited it) and the same origin spelled more than one way.

When somebody tries to pay and it does not work

Sarah Griffith tried to pay \$100 for two glow seats at **10:47pm on 12 September** and her bank declined it for insufficient funds. Her row read **new / unpaid** — true, and exactly what it read before she ever opened the link. The webhook only listened for success, so the attempt existed nowhere but in Stripe's own dashboard. It was found because you went and looked.

A guest who tried and was stopped by her bank needs "*your seats are still there, try another card*". A guest who never opened the link needs a nudge. Those are opposite messages, and until now both looked the same on screen.

So a declined payment is now recorded on the booking and shows in the Bookings list as **⚠ TRIED TO PAY — DECLINED**, with the reason and the time on hover. You also get one email naming what to say.

Nothing is sent to the guest automatically. Stripe already told them on the spot, in their bank's own words, and an automatic second message from us an hour later only causes the phone call. What you get instead is a **Text** button on the row, which opens your own messaging app with the words already written — you send it, from your phone, when you choose to.

That text leads with *your spot is still held* and **never says why the card was declined**. We know the reason and the guest does not need it from us; "your bank said insufficient funds" is a humiliating thing to receive from a boat company, and it is their bank's business, not ours.

Every guest text now ends by saying a real person is on the other end and they can reply — because they can, and a number that looks automated gets no answer when you actually need one.

Nothing is cancelled and no seat is released. A Stripe checkout link stays open for 24 hours, so most of the time the link they already have still works — the email tells you whether it does, by its real expiry rather than by assumption. Cancelling on a decline would turn a retryable moment into a lost booking.

It says when the fault is ours. An expired API key or an amount Stripe will not take is not something a guest can retry their way out of, and that email arrives as **PAYMENT BROKEN OUR END** instead. An unrecognised decline code is passed through in Stripe's own words and flagged as not understood rather than being softened into "their card didn't work" — which could otherwise hide a fault of ours while somebody sits there unable to pay.

The badge clears itself the moment a payment succeeds, so a paid booking never wears an old decline.

A charter that was paid for and never happened

Weather, a breakdown, a guest who cannot make it. They have paid, the day is gone, and no new date has been agreed. That booking is **owed** — a status added on 5 September 2026 because there was no honest word for it.

It is not *cancelled*. Cancelled means money went out or is going out and the relationship is finished. Owed means the opposite: their money is still here, they still want to go, and the business owes them a boat. Christian Gehring sat marked cancelled for two months for want of this distinction, and nothing ever put him on a list.

What the status changes:

- **It does not hold the day.** An owed charter has no date, so it blocks nothing on the availability calendar. This is the one status where a real booking with real money behind it occupies no day at all.
- **It counts as active.** It shows under the Active filter on the bookings table, in amber, because it is work outstanding rather than history.
- **It is not a cancellation** in any count, so conversion figures stop being wrong in the business's favour.
- **Pearl reports it every morning** until it is settled. If there is no phone or email on the row she raises it as urgent, because a charter that cannot be rescheduled is money owed forever.
- **Money can still be attached to it** in the ledger. The payment is real.

When they pick a date, set it back to **booked**. Owed is a waiting room, not a destination. The standing rule when you contact them is to offer a weekend, not a refund — if they want the money back they will ask.

Where to see them: Overview → Guests → "Charters we owe"

Each one shows the guest, what you are holding, how long it has been, and which boat. Three buttons:

- **Text it** — opens your phone's messaging app with the message already written. On a desktop it says "*phone only*" and refuses, because a desktop has nothing to hand an `sms:` link to and a message you think you sent is worse than one you know you didn't.
- **Preview** — shows the exact wording, with **Copy** underneath for sending it another way.
- **Email** — only appears when there is an email on file.

The wording is not editable, on purpose. None of the drafts mention a refund and none of them apologise at length, because opening with "*do you want your money back?*" invites the answer that ends the relationship when what the guest wanted was to go boating.

The section is **not there at all** when nobody is owed. That is the normal state.

When a charter's money can't be found

The Overview also says "*N income rows are not tied to a charter.*" That is not the same as money going missing, and the difference matters more than the number:

- **Unlinked** — the money is on the books, but nothing joins it to the charter, so anything asked from the booking's side answers "no money". **Fix the link. Do not add a row.**
- **Missing** — nothing in the ledger matches it at all. Find out what happened, then write one.

Adding an income row for money that was already recorded **doubles it on your tax report**. Christian Gehring's \$520 looked missing for exactly this reason: the Zelle row had been there since 9 June and simply wasn't tied to his booking.

Pearl's morning check now reports this, and says which of the two it is rather than leaving you to guess. Some unlinked income is perfectly correct and always will be — the May 2026 Glow Party seats are real income with no booking to attach to.

Booking numbers never change

Every booking is `NY-YYYYMMDD-NN`. The date inside it is the date the charter was **first booked for**, not where it ended up, and the `NN` is just the order it was taken that day.

So when a charter is rescheduled, **the booking number stays exactly as it is** and the date column moves instead. That is deliberate: the number is what a guest quotes on the phone, what the ledger points at, and what an old email says. A number that moves is a number that stops matching the paperwork.

There was one exception, made once, on 5 September 2026: `NY-20260711-GEHRING` was written by hand and ended in a surname instead of a number. It became `NY-20260711-02`, with its ledger entries moved in the same transaction. That was a one-off correction to an id that never conformed — not a precedent.

Availability

Block days per vessel. A day with bookings that do not fill it shows as partially booked, calculated from the summed hours of that day's charters.

A partly-booked day now says when it is taken. A guest complained on 8 September 2026 that clicking a date told him it was partially booked and nothing more — which hid the one fact that decides whether he can still come. The day now reads its windows: `7-11pm`, `10am-2pm`.

The information was always there; the calendar was throwing it away and keeping only the total hours, because all it needed to answer was whether the day was full.

A booking with no start time on record is left out rather than guessed at. Twelve of the forty-two are like that — mostly older ones taken by text — and an invented window would be worse than a vague one, because a wrong time is something a guest will act on.

A booking you confirm by text now blocks its date. Until 5 September 2026 the public calendar only knew about charters paid for by card, because the Stripe webhook was the only thing that wrote a diary entry.

Anything you took over the phone and marked **booked** left the day still on sale — you could have sold it twice. Marking an inquiry booked is now enough.

Nothing else changed: an inquiry still blocks nothing, and a charter paid by card is counted once rather than twice even though it exists in both lists.

Money

Income & expenses

Every dollar in or out. The form on the left adds an entry: Income or Expense, a category, an amount, a date, and for income the origin it came from. Linking an entry to a booking makes it appear in that booking's profit.

Below the list: breakdowns by category, profit per booking, and commission lost to the platforms.

What is actually in the bank

The ledger answers *did we make money*. It does not answer *will Thursday's bill clear*, and those are genuinely different questions — on 31 August 2026 the season was in profit while the business account sat at **minus \$11.36** with two rejected Optimum payments against it.

So the Money card on the Overview now opens with the account balance.

It is a reading, not a live figure. Nothing here is connected to Woodforest. Somebody reads the app and records what it said, with the date it said it. The age is shown beside the number for exactly that reason — today , 1d ago — and past a week it turns amber, because a stale balance read as current is worse than no balance at all. The colour of the figure itself is the money: green, amber below \$250, red if it is negative.

Readings are never edited. A figure that was wrong is corrected by recording a newer one; the old reading stays. That is the whole reason for keeping readings rather than one number that gets overwritten — the 31 August overdraft is only visible because it was kept, and a balance history that begins at a healthy number hides the month that went wrong.

Two are on file, both taken from the same screenshot on 16 September 2026:

As of	Balance	
16 September 2026	\$934.30	after the \$824.70 transfer in on the 14th, and the \$45.32 Optimum payment out the same day
31 August 2026	– \$11.36	overdrawn. A \$12.00 service charge took it under; the Optimum payment had been rejected twice, on 12 and 18 August, and did not clear until 14 September

Wells Fargo is being retired

Your decision, 16 September 2026: **the business pays from Woodforest now.**

Filing an expense against *Wells Fargo Statement* dated on or after that day raises a warning under the origin dropdown. It does **not** stop you. If the money really did come out of Wells Fargo, file it truthfully — the warning's job is to remind you to go and move whatever is still charging that account.

Rows dated before the cutover are ordinary history and say nothing. There are 171 of them going back to June 2025 and every one is correct.

The morning check watches the same thing from the other side. The warning only covers rows you file by hand. A statement import, an agent recording a cost, or an autopay you have forgotten is still pointed at the old account never touches that form — and those are the ones that matter, because nobody chose them. The check reports and does not block: moving direct debits takes weeks, and a checker that shouted about a cable bill would be turned off by Thursday.

As of 16 September nothing has been paid from Wells Fargo since the decision. The newest row on it is 25 August.

Reconciliation

Answers one question per booking — is this charter's money actually on the books? Matching is on the real foreign key between a booking and its ledger rows, not on date and amount.

Boatsetter pays in two legs, the boat and the captain fee, often days apart. Two income rows against one charter is normal, not a duplicate.

Tax Report

Pick a year for totals, a CSV export, and breakdowns.

Money held for a charter that never ran is not counted as income, in this report or in "Season in" on the Overview. It is a deposit against a trip that has not happened — if the guest asks for it back, it goes back — so it becomes income in the year the charter actually sails.

It is never just removed. Both places name the excluded amount, because your bank statement will show that money arriving and the two figures have to be reconcilable. The Overview shows it as "**Held, not earned**"; the Tax Report shows an amber note above the totals saying how much and why. Anyone marked **Owed** in Bookings is where it comes from.

Income is split **by vessel** and **by origin** rather than by category — every reservation is logged under the single category "Reservation", so a by-category panel would be one row totalling everything.

It also shows **average per charter** and **average per hour**, which are the two numbers pricing actually turns on. These come from bookings rather than ledger rows, because a Boatsetter charter produces two income rows and counting rows would halve the apparent value of a trip.

If income has no vessel recorded, an amber note says how much. That is a prompt to fill it in, not a rounding error.

Gift certificates

Certificates bought from the public site, and what has been redeemed against them. The tab sits under **Money** because that is what one is: money taken now for a charter owed later.

Each shows its code, who bought it, the face value, what is left on it, and its state. A certificate that has been paid for and not yet redeemed is a **liability** — the money is in the account but the trip has not been given. It is the same shape as a charter someone paid for and never took, and it should be read the same way.

To redeem one, apply its code at checkout or against a booking. Partial redemptions leave the remainder on the certificate.

They do not expire on their own. Nothing in this system voids one for age, so an old certificate is still owed unless you decide otherwise — and that is a decision to make deliberately, not to discover when someone turns up with it.

Subscriptions & bills

Every recurring cost the business carries, normalised to a monthly figure so weekly, monthly and yearly items can be summed.

It used to hold software and boat storage only. On 13 September 2026 the household accounts moved in beside them — electricity, gas, water, sewer, trash, phone, internet, storage — because the office is at home and a bill you never recorded is a deduction you never claimed. Bringing them in recovered **\$2,451.56** of payments that had been made and never entered anywhere.

The four columns that decide what a bill is worth

Where — which address it belongs to. The move in April 2026 means most services have one account at each, and several of the old ones ran on for months afterwards. Two live accounts with the same supplier usually means the old one was never closed, and the tab says so out loud when it sees one.

Amount — and **blank is not zero**. Blank means nobody has found out what it costs yet; `0` means it is confirmed free. The two look identical in a total and mean opposite things, so they are stored differently and shown differently: a blank amount is listed by name under the totals as an open question, and adds nothing until somebody answers it. Supabase and Vercel really are free. Microsoft and ElevenLabs simply have not been looked up.

Business % — how much of that bill is the business's, and **blank means nobody has decided**. It is not zero. A bill with no share set stays out of the deductible total entirely rather than being guessed at in either direction.

Ended — the date the account closed. Setting it is what closes an account; `Active` follows it automatically, so a closed bill can never keep counting toward the monthly total because two fields disagreed.

The two totals, and why they differ

The first is what leaves the bank every month. The second is **the business's share of it** — the only figure that belongs on a tax return.

Boat storage and software are 100%. Whole-home utilities are a proportion: the office is a 10 × 10 room in a 1,900 sq ft home, which by floor area is **5.26%**. You chose to claim **10%**. Both numbers are written into every bill's note on purpose, because a percentage in a tax record is only worth the basis somebody can point at a year later, and an accountant should see the measurement as well as the figure claimed.

Three bills are not apportioned by floor area, because floor area is the wrong test for them: T-Mobile (a business account in the LLC's name, 13 lines), Optimum and AT&T. Those are set to 100% on your instruction. The note on each records that the 100% is your account of how the line is used and is not something the bill itself establishes — which is exactly what an accountant will ask about first.

Personal subscriptions, and why they are here at all

Netflix, Hulu, Spotify, Xbox Game Pass, the car note, renters insurance. You wanted one place to see everything you pay for monthly so duplicates and waste could be found across the lot — and the very first look found two identical Google One subscriptions on two different Google accounts.

They are counted nowhere near the business. Not at 0%, which would still put Netflix inside the monthly cost and merely leave it out of the deductible half. The split happens before any arithmetic: personal rows have their own line with their own total, they never enter the business figures, and they are stripped out of the subscriptions CSV that the Tax Report exports for a bookkeeper.

The **Side** button on each row moves a bill between the two. It is a button rather than a fixed property because the answer can genuinely change — Meta One Advanced looked personal until its billing screen showed it subscribed to The Nauti Yacht page, at which point it was a marketing cost.

Grouped, and folded shut

Thirty-odd rows in one run is a wall, not a list. Bills sit under their category — Utilities, Storage, Software, Hosting, Other, then Personal last — and every heading carries that group's own count, monthly total, business share, how many have no amount yet and how many are closed.

Groups start closed. Click a heading to open it. Inside a group the biggest bill sorts first, because that is the one worth arguing about, and anything unpriced sinks to the bottom where it reads as a question rather than as a cheap item.

The calendar

A total says how much. It does not say how much *this week*, which is the question you actually ask when deciding whether something can be paid now. The calendar draws the month, with each bill on the day it lands.

It will not pretend. Only the bills whose due date came off a receipt are on it; the rest are named underneath as bills it cannot place. An empty Tuesday has to mean nothing is due on Tuesday, not that nobody has looked. A monthly bill repeats on its day and is drawn every month; a yearly one appears only in the month it actually falls, because spreading Peacock across twelve squares would turn one \$79.99 charge into an imagined \$960.

It earns its place immediately: Blotato, Claude Max and ElevenLabs all fall due on the 28th — \$159.53 in one morning, which no list sorted by name would show you.

On a phone

The table stacks into one card per bill, each row labelled by the column it lost. Nothing scrolls sideways, and the calendar scrolls inside its own box rather than widening the page.

Marketing

Media

The public gallery, grouped by package — bachelor, birthday, corporate, glowz, night, partycove, tubing. Grouping makes it obvious which package is thin.

Captions edit in place. + **Add** adds a tile to that category.

New images belong in the site's `public/gallery/` folder and are referenced as `/gallery/name.jpg`. They are then served from our own repository. Most of the older tiles still point at BrandCrowd, a logo-design service — if that account ever lapses, those images disappear, which is why new ones go in our own folder.

Media Drafts

Every post Coral has drafted, laid out as cards and boxed by the day it goes out. Two days side by side, soonest first.

On a card

- The photo or clip it goes out with, or **No media attached** with a link to add one. Instagram and TikTok refuse a post without one, so a scheduled card showing that warning will fail on the day.
- The caption in full, so you are approving the words against the picture.
- The platform and when it goes out.
- **I posted it myself · Don't post**

Posted, denied and past drafts sit in a collapsed group above the days — a record, not a to-do list.

What goes out on its own, and what does not

A post marked SCHEDULED will publish by itself. Siren runs each morning and puts out whatever is due that day. Giving a draft a date is the permission to post it, so she does not ask again — there is no second confirmation, and no message the morning it happens.

Approving a post that already has a date is the same permission. Siren puts those on the schedule herself at the start of each run, then publishes them when their time comes. You do not have to press anything a second time.

That was not always true, and it cost a post. On 12 September 2026 the Facebook copy of a Boatz & Glowz post went out at its 10:30 slot and the Instagram and TikTok copies of the same post did not — identical date,

identical time, and those two were sitting at APPROVED rather than SCHEDULED. Nothing errored and nothing logged it; it was spotted by looking at the feed.

Permission still stops there. Nothing else is ever published: not a draft Coral has proposed, not one you have approved **without a date**, and not one you have rejected.

So the queue is safe to leave alone **until a post's date arrives**. Up to that morning you have as long as you like to read it. Once the date is today, the next thing to touch it is Siren.

To stop one, press **Don't post**. That is the only thing that takes it out of her way.

Saying what is wrong with a post

Discuss is on every card, at every stage — including after it is scheduled, which is exactly when "not that clip" tends to get noticed. It replaced **Preview** on 8 September 2026, which showed the caption and the media a second time underneath the ones already on the card, and so cost a click to learn nothing.

Pick what is wrong from the list, then say whether the post survives it:

Reason	Means
Wrong photo or clip	It does not match the post — wrong guest, wrong trip, wrong thing happening
Find a better one	Right footage, weak shot. There is better in the same folder
Caption needs work	Wording, tone, hashtags or a wrong detail
Wrong day or time	Right post, wrong slot
Too similar to another post	Same clip, angle or message as something already queued
Guest or privacy problem	Someone in it should not be, or it names the wrong guest
Not right for us	Wrong message for the business, whatever the media
Something else	Say what in the note

Wrong photo or clip and **Find a better one** are deliberately separate. Wrong means find the *right* one, and sends whoever picks the replacement to a different folder; better means find a *stronger* one, usually from the same shoot. Both offer **Swap the media now** so you can replace it without leaving the card.

Then either:

- **Keep it — send back for changes.** The post stays in the queue and moves to *Needs work*.
- **Don't post it at all.** It moves to *Rejected*.

Neither of these deletes anything. Rejected is not gone: **Back to review** brings it straight back. Deleting is a separate button that only appears on a rejected card, behind a confirm. The note is optional — "wrong clip" is already a complete answer.

Why the reason is a fixed list

A rejected post used to record nothing at all. It vanished from the queue and the reason lived only in your head, so the same mistake could be proposed again the following week and nothing could count how often that happened.

The list is short and fixed because the point is counting. *"Six of the last ten were killed for the wrong clip"* is a fact that changes what Coral does next week; ten sentences roughly saying that are not. The free-text note still exists — the two are written together, and both show on the card afterwards, including on rejected posts, so a post you killed still says why when you come back to it.

"I posted it myself"

This button used to say **Mark posted**, which did not explain when to use it. For anything Siren publishes you never touch it — she marks her own as posted and records the live URL. It is only for posts that go out by hand, which is still most of TikTok.

What each account will actually take

	photo	video
Facebook	yes	yes — publishes as a Reel
Instagram	no	yes
TikTok	no	yes

A still cannot be published to Instagram or TikTok at all. The publisher reaches Instagram only as reels and stories, and both need video. Siren treats an Instagram or TikTok draft with a still as blocked: she will not attempt it, and reports it instead.

So a photo-only post can go to Facebook, and if you want it on the other two it has to be posted by hand from your phone. Facebook takes either, freely.

Every Instagram post tags your personal account

Instagram posts carry a **Collaborator** tag on `austinhefty` and `brookeashley_05_`. A collab is not a repost: Instagram puts the one post on all three profiles' feeds and grids, and the likes and comments pool rather than splitting across separate copies.

You each get a notification to accept, once per post. Nothing fails if you don't — the post publishes to the business account regardless and simply doesn't appear on your profile.

Instagram caps this at **three** collaborators, so there is room for one more. The list lives in `lib/socialPosting.js`, not in an agent's brief, so it cannot be forgotten on a run.

There is no equivalent on the other platforms, and there cannot be. Meta removed the ability for apps to post to personal Facebook profiles in 2018, and TikTok's API writes only to the authorised account. Snapchat has no organic posting API at all and is not a platform the publisher supports. On those three, sharing to your own Story is a manual tap — and worth doing, because your Story reaches local friends who might book or refer.

Five hashtags, never six

Blotato rejects an Instagram post with more than five: "*Instagram allows a maximum of 5 hashtags per post.*" It does not trim, warn or retry — the draft stays marked **scheduled** and looks perfectly healthy.

The sober-captains video failed this way on 4 September 2026 and again on the 5th with the caption untouched. On the 8th an audit found **seventeen of twenty-one** Instagram drafts over the limit at once — the better part of a fortnight of Instagram silently not posting, with nothing to say so until each day came and went.

The limit is now applied to every platform, because nothing this business has ever published carried more than five on any account, and a dropped generic hashtag costs less than a post that never goes out. `drafts.js --check` reports any caption over five as **T1**.

Where a clip came from

Every draft cut since 8 September 2026 records its source in the photo hint — which charter folder, which file, which second of it.

That exists because of a real confusion: a night-cruise clip was read as coming from one guest's folder when it was another's, and nothing on the card said otherwise. Older drafts whose media predates this are left blank rather than guessed at.

Approved with no date

An approved post with no date **will never go out**, and that is the one gap still left on purpose. Approved with a date means "yes, and then" — Siren schedules those herself. Approved with NO date means you have said yes but not when, and nothing invents a date for you.

Coral now proposes a date for each of these in her daily status, and raises one as a board item if it has been waiting more than three days. The Overview also flags them under **Needs attention**. Set the date with **Reschedule**.

Comments

The queue of comments nobody has answered, from Facebook and Instagram.

It exists because of one day. On 6 September 2026 a Boatz & Glowz post drew four comments challenging the operation — liability, life jackets, drink-driving, litter on the shoreline — and the first sat **twenty-three hours** before anyone saw it. The people reading a thread like that are the ones who never comment, and an unanswered accusation reads as conceded.

A thread is colour-coded by how long it has waited: **fresh** under six hours, **waiting** past six, **overdue** past twenty-four.

Nothing here posts on its own. Siren writes a suggested reply and it sits in the box; you can send it, edit it, or empty the box and write your own. A reply is public, immediate and attributed to the business — every other agent in this system proposes and you decide, and a comment thread composed in response to something hostile is the last place to break that.

TikTok comments are not here, because the publishing API does not expose them. They remain a manual job in the TikTok app. Saying otherwise would leave you believing a channel was covered while nothing was watching it.

Messages

The direct-message inbox, from Facebook and Instagram.

Nothing read these before it existed. There was a comments tab and no messages tab, and not one of the scheduled crew had DMs in their brief — so a message arrived, sat there, and there was no point at which anybody found out.

The first look at the real data found two unread messages from 7 September, a day and a half old, from the man who had posted publicly about litter on the shoreline. He was offering two of his own boats to help clean up after the glow party.

It is built to the same shape as Comments on purpose. They answer the same question on different channels, and two panels that disagreed about what "waiting" means would be worse than one.

This is no longer true of sending, and that changed on 17 September 2026. Nothing YOU see here sends on its own — but most DMs are answered by an automatic reply seconds after they arrive, before this tab is ever opened. See *Messages answer themselves* below.

"I answered this elsewhere"

The console cannot see a reply you typed in the Facebook app. Blotato only knows about messages Blotato sent, so when you answer a DM from your phone — which is most of them — that reply never comes back to us. The thread keeps sitting there in red saying **never answered**, and it is wrong.

It was wrong about two threads on 17 September 2026, one of them for ten days: the man offering two of his own boats for the shoreline cleanup had been answered the same week, and the console had no way to know.

So each waiting thread carries a small **I answered this elsewhere** link under the reply box. Press it and the thread goes quiet: the red flag clears, the colour drops, and it stops counting toward the number on the Messages tab. The card then reads *answered by you*, with **put it back** if you pressed it on the wrong one.

It sends nothing. It writes a note against that thread in our own database and touches no platform.

If they write again, it re-opens on its own. The mark records *when* you dealt with it, not a permanent "done" — so anything they say afterwards puts the thread straight back in the queue with its flag. You never have to remember to un-mark it.

And it has no effect whatsoever on the automatic reply. That fires inside Blotato the moment a message arrives and never reads our database, so their next question gets an instant answer whether or not you marked the thread. Nothing on this screen can switch that off — that switch is in Blotato, on automations 5989 and 5988.

Messages answer themselves. Comments do not.

Your decision, 17 September 2026, and it is the one place in this system where something goes out without you reading it first — so it is worth knowing exactly how far that goes.

The difference is the size of the audience. A comment is published under the business's name to everyone reading the thread. A DM goes to one person and can be corrected in the next sentence. Speed is worth more in an inbox; caution is worth more in a thread.

No agent sends any of it. The automatic reply is a Blotato automation, which is a platform feature rather than a crew member. Nothing Siren or anyone else runs sends a message to a person — that rule is intact.

What actually goes out. Within seconds of a DM arriving, on any platform, at any hour:

> Thanks for the message 🙌

>

> Easiest way to get you sorted is a text — drop your number here and we'll send

> you dates, prices and a link to lock it in.

>

> Or have a look at the site below. A real person reads these, so ask away.

>

> — The Nauti Yacht

It asks for the number because a number turns a DM into a text thread, and the pay link goes by text. It is a fixed string: **it never quotes a price, a date or a seat count**, because it cannot know any of them.

It fires on keywords, and the list was nearly useless. It held thirteen words — book, booking, price, prices, pricing, cost, available, availability, rent, rental, glow, tube, tubing. On 17 September a guest wrote "*Hey bub, you got any seats open for Saturday?*" and matched **none of them**: the single most common question this business gets, missed by one word. He sat for two hours and was found by eye, not by the system. It now holds forty, including seat, spot, room, space, open, the days of the week, "how much" and "what time".

The hazard, and what watches it. A keyword match is blind. "*Your prices are a scam and I want a refund*" contains `price`, so the machine sends that person the packages link. Nothing stops it — the automation cannot read.

So every waiting thread is now checked against a set of rules after the fact, and the Messages tab flags any the machine answered that it should not have, in red: **"Answered automatically, and should not have been"**, with the reason. Read what went out before you reply to one of those.

What is never the machine's to answer — these always come to you, and they are the reasons you will see on a card:

Held because	Meaning
safety	someone may have been hurt, or something was damaged
legal	lawyers, insurance, liability, a refund or a chargeback
complaint	it reads as a grievance
money already moved	it asks to change, refund or reschedule something paid for
solicitation	a pitch, not a guest
NDA	it touches the Lake Bryan charters
photographs	somebody's picture, or a request to take one down

Plus three structural ones: anything over 320 characters, anything asking more than two questions, and anything matching nothing the system can answer from its own data. That last is the important default — **an unrecognised message is not a safe message, it is a message nobody has understood.**

A price question wrapped in a complaint is treated as a complaint.

Siren drafts the held ones. Her reply appears in the box the same way it does on Comments, with her name on it, and you send it or type over it. If they write again before you get to it, the draft is marked stale and will not pre-fill — an answer to a question that has been overtaken is worse than an empty box.

Asking for Google reviews

The panel at the top of **Testimonials** lists everyone who could be asked, with the freshest charter first.

Four filters: **Still to ask**, **Already asked**, **Not asking**, and **Every charter** — which means every charter, archived ones included.

Button	What it does
Text it (green)	Opens your messaging app with the message already written, and ticks them off. Phone only — a desktop has nothing to hand an <code>sms:</code> link to.
Preview (blue)	Shows the wording first, and can copy it for pasting into a platform message thread.
Don't ask (amber)	Takes them off the list for good.

Don't ask is for the guest you will never ask — one who damaged something, or who you would rather not hear from. They move to **Not asking**, come out of the counts, and can be restored with **Put back**. Nothing is deleted. Before this existed the only way to clear such a guest was to mark them asked, which puts something untrue in the record and then hides it behind a tick.

Ask via only appears for guests with no phone number. Those cannot be texted at all, so the platform thread is the only way to reach them, and the column is the instruction. A guest with a number just says *Text*.

Two rows say **extra contact** instead of a charter reference. Those are people who sailed on somebody else's booking and whose number you kept — see *Extra guest contacts* under Bookings. They can be asked like anyone else.

Testimonials

Reviews shown on the public site, plus the panel for **asking past guests for a Google review**.

Neither booking platform hands over contact details, so every number was typed in by hand — a guest with no number is a review that never gets asked for.

Text it opens your phone's messaging app with the message already written and ticks the charter off. **It does nothing on a desktop** — this has to be done from a phone.

Photo Requests

Guests who asked for their photos and have not been sent them. The badge counts what is still **owed**, not the total, so once the queue is clear it says nothing.

Longest wait first, deliberately not newest first. The guest at the top is the one most likely to have given up on us. The days-waiting figure beside each name is counted in whole days from the day they asked.

The link cannot be filled in from here. The files are on your machine, not in this database, so the button opens the message with the greeting already written and leaves you to paste the album link — the whole job reduced to one tap and one paste. Marking it sent is what clears it from the queue.

Setup

Packages & pricing

Per-package prices, per-guest rates, and the hourly grid per vessel for weekday and weekend. Every change is logged to price history.

Add-ons · Coupons

Extras that can be attached at checkout, and discount codes with optional expiry and usage limits.

What comes free with what, and what is charged:

Add-on	Price	Free with
Balloon Package	\$40	Birthday Party
Champagne on Ice	\$25	Bachelor / Bachelorette
Full Decoration Package	\$60	— <i>(contains the balloons and the champagne)</i>
Grill Service	\$25	Night Cruise

The **Full Decoration Package is the other two plus table setup and themed decor**. Bought separately the balloons and champagne come to \$65, so the bundle saves \$5 — and a guest who ticks all three is charged \$60, not \$125. That is enforced in code, not by remembering.

Grill Service is the cooking only — the guest brings the food. It is a small electric grill, about two steaks at a time or four burgers at a time, and the boat has to be at anchor, so allow twenty to thirty minutes. It can be added to any charter; anchoring off mid-trip is normal.

A charter with no stated occasion is a Tubing / Wakeboarding charter. That is the default, not a guess.

These rules live in `lib/addons.js` and are the single source for the booking form, the price quoted and the FAQ. Change them there, not in the copy, and run `node scripts/test-addons.js` afterwards — 51 assertions cover the pricing.

The top bar

Button	What it is
The version badge	The version of the whole system, top left beside OWNER CONSOLE. The same number appears under the Owner console button on the public site. Both read it from <code>package.json</code> , so they cannot disagree with each other. See Versions and going back below.
 what Pearl last said	Appears only when she has said something. The words are shown; the speaker button reads them aloud in her voice, and pressing it again stops her. The x dismisses that message — the next one brings the strip back. There used to be an Enable Pearl button here whose only job was to satisfy the browser's rule that audio needs a click first. Any click now does that, so it is gone.
 On the dock	The phone page: ask for reviews, log engine hours. Everything on it needs a phone — a text link has nothing to open on a desktop, and an hour reading is taken at the boat. Worth adding to your home screen.
 Manual	This document.
 Back to site	The public website.
Log out	Ends the session.

The waiver draft link is gone. The Release and Waiver is published in the **Terms & Conditions** tab on the public site, and booking now records that a guest accepted it.

A booking that came in by text

Most bookings do not arrive through the website. They arrive on WhatsApp, by text, or on the phone — and that channel is the biggest and the worst tracked. The 8 Aug 2025 charter was never written down at all and had to be reconstructed from photographs a year later.

This is the route that stops it happening again.

1. **Send the message to Claude** — a screenshot is enough. 2. **It gets logged straight away** as an inquiry with a real reference (NY-YYYYMMDD-NN), before anything is promised. Availability, boat capacity and the price are checked against the live data, not from memory. 3. **You get a reply to send**, written for what that guest actually asked. 4. **If they say yes**, you get a link: `thenautiyacht.com/pay/<id>`. It shows them the charter on our own site and hands off to Stripe from there. Never paste a raw `checkout.stripe.com` link into a text — it looks exactly like a scam, and the more careful the customer the less likely they are to tap it. 5. **When they pay, the rest happens on its own**: the booking is marked paid, the calendar closes that day for that boat, and a confirmation goes to the guest with you copied in. A website inquiry also gets a booking row created at this point; a booking taken by text already is one, so nothing is duplicated.

Doing it yourself, from the phone

From v2.4.0 you do not need to ask for the link. Open **Bookings** on your phone and each row carries the text it needs:

- **Text payment link** — a lead with a price on it. Sends them their own `/pay/<id>`, already filled in.
- **Text to confirm** — a lead with no price yet. Asks them to confirm so you can price it and send the link after.
- **Text reminder** — a booking that is already paid. Day, departure time and where to meet. Boatz & Glowz says Scott's Ridge by name; every other charter meets wherever that one was arranged, so it stays general.
- **Text about owed** — a charter paid for that never sailed.

These only work on a phone. A desktop has nothing to hand an `sms:` link to, so it tells you that rather than appearing to send.

A lead only gets a payment link if its row has a package and a price. That is what the link charges for. A booking logged by text with neither falls back to *Text to confirm* until you add them.

The one gap to know about. Between "yes" and "paid", nothing holds the date. The calendar only closes on a real booking, so a second party could still take that boat. If the charter is soon or the date is in demand, block it by hand in **Bookings** → **Availability** and unblock it if they never pay.

Capacity counts the captain

The vessel capacities INCLUDE the captain. The Nauti Yacht's 12 seats **11 guests**. A party of 12 needs the Explorer. This is the easiest thing on the whole site to get wrong when quoting quickly.

If a payment and the console ever disagree

A guest says they paid and the booking still reads unpaid — ask Claude to resync it. That re-reads Stripe's own record and repairs whatever is missing. It only ever copies from Stripe, so it cannot invent a payment that did not happen.

Versions and going back

The number beside OWNER CONSOLE says which release you are looking at. It exists for one moment: something is wrong and you need to know what to go back to.

What earns a new number

The crew doing their job is not a revision. Coral drafting posts, Siren publishing, Pearl filing a brief, a booking landing, a photograph being tagged — none of that moves the version. That is the system running, not changing.

Size	Moves	What it means
Small	1.5.0 → 1.5.1	Formatting, wording, a colour, a label. Nothing behaves differently.
Moderate	1.5.1 → 1.6.0	The logic of an agent, or a workflow. A status meaning something new.
Large	1.6.0 → 2.0.0	A whole layout changed, or a completely new large-scale idea.

Three numbers, always — `1.5.1`, never `1.5`. Unsure between two sizes? **Take the smaller one.** A small change that turns out to have been moderate costs nothing; a moderate one recorded as large makes the history lie about how much moved, and the history is the entire point.

Nobody cuts a release without asking you

A revision draws a line under a batch of work, and only you know where the batch ends. The rule is to finish the work, say what changed, say what size it looks like, and **ask** — because it was broken twice on the day it was written.

This manual is re-read at every release. Two checks stand behind that. One compares the tab names in it against the tabs the console actually has, so a rename can't quietly leave this document describing a screen that no longer exists. The other records which version the manual was last read against, and complains as soon as the code moves past it.

The second one exists because the first cannot do everything. On 13 September this manual said "*The guest gets nothing*" about a declined payment. That was true when it was written at one in the morning and false by six that evening, when the Text button arrived — and no checker reads English for truth. A sentence that quietly stops being true is caught by a person re-reading it, and the release is the moment anyone reliably does.

What you actually get

One `.zip` in `AI & Website/releases/`, and only ever the newest one: the crew briefs, the protocol, the schedules, the permission rules, the hand-written skills and the shared scripts — everything git does not version. Older *release folders* are deleted; **every git tag is kept forever**, so the code for each version is still on GitHub.

The skills were added on 14 September 2026 and had been missing the whole time. They live in `.claude\skills`, a different folder from the crew's `.claude\scheduled-tasks`, so nothing was picking them up — which meant `/watch`, the tool that turns a video into timestamped frames and a transcript, would simply not have existed after a restore, despite the whole media pipeline being built on it. Skills synced down from `claude.ai` are skipped: those are somebody else's copy and come back on their own.

Where the files live, and why it is split

Two places, and the split is about what Google Drive can back up.

	Where	Backed up by
Business files — Photos, Finance, Legal, releases, this manual	<code>Documents\MyFiles_The Nauti Yachti LLC</code>	Google Drive , which mirrors that folder
The app — code, dependencies, build output	<code>Documents\Nauti-yachti-app</code>	GitHub , off-machine and stronger for code

Drive cannot sync a working tree. Until 15 September 2026 the app lived inside the mirrored folder, and Drive had quietly given up on the whole thing and marked it with a red X — 1.5 GB of build output and 755 MB of dependencies were more than it would take. Nothing announced it. The folder looked normal while nothing in it was being backed up, and it holds 53 GB of charter footage.

So the app moved to the local disk. **This is not a gap in your backups.** The code is on GitHub, tagged at every version, which is better protection than Drive gives it. Drive still holds everything git does not version — the photo library, the finance records, the release zips, this manual.

A checker enforces it now. `check-consistency.js` walks the business folder for `.git`, `node_modules`, `.next` and `.turbo`, and refuses to let a release through if it finds any. It is deliberately about backup coverage rather than correctness: the point is that nothing can drift back in unnoticed. It found three stale pointers in old distributable test builds the day it was written.

One consequence worth knowing: a file that git ignores AND that lives on the local disk is in no backup at all. There is one — `conf\token.txt`. Its contents were exposed in a working session on 15 September 2026, and the file now holds a note saying so rather than a token. **That credential still needs rotating.** Nothing in the app or

the crew scripts reads the file, so nothing is broken by it sitting there — which is exactly why it would be forgotten.

`CHANGELOG.md` carries one entry per version, written to be read on a bad day. If an entry cannot tell you whether to restore that version, it was not written properly.

The full reasoning, including why this is deliberately **not** SemVer, is in `AI & Website/VERSIONING.md`.

Boat

Maintenance

Service items per vessel against elapsed months or engine hours, plus engine-hour readings and fuel fill-ups. Items flip to overdue automatically.

Nothing here can judge anything until engine hours are logged. Thirteen items are configured against hour intervals and the panel will keep reporting "nothing overdue" while it has no readings — which looks identical to a healthy fleet. Log a reading after each outing from **On the dock** (see below); the intervals then work on their own.

Logging fuel writes the expense for you — the one other place besides completed bookings where the ledger fills itself in.

What runs on its own

Nine scheduled tasks: eight named crew members and the standup that files their status. **None of them post, spend or send anything to a guest without you.** Only one acts outside the business at all.

They run in dependency order, so the summary comes after the things it summarises:

When	Who	What she does
Daily 8:19	Nauti Penny · Accounts Receivable	Money in. Payouts against the ledger, and anything paid for that never ran.
Daily 8:38	Nauti Coral · Content Producer	Drafts a post from real fleet photos, audits the queue, and checks what Siren actually published.
Daily 9:04	Nauti Siren · Publishing & Brand Safety	The last gate before anything is public, then publishes what passes. Drafts replies for unanswered comments and for the DMs the auto-reply held back.
Mon 9:26	Nauti Joy · Guest Relations	Who to ask for a review, and who was left hanging without one.
Mon 9:41	Nauti Reef · Revenue Growth	Money the business is not collecting.
Mon 10:04	Nauti Shelly · Accounts Payable	What is being paid for against what is actually used.
Mon 10:29	Nauti Nova · Market Research	The outside world. Reports nothing most weeks, by design.
Daily 10:51	Crew Standup	Files a one-screen status for all eight.
Daily 11:18	Nauti Pearl · Chief of Staff	Reads everything, decides what reaches you.

The email the site sends

Four things a guest can do on the website send mail, and **every one of them now emails both the guest and you**. Three of the four used to be one-sided: an inquiry and a crew-list signup told only you, and a gift certificate told only the recipient — so a guest who typed their details into the form got silence, and a certificate could be bought without the business hearing about it.

What the guest did	They get	You get
Sent an inquiry	acknowledgement, "we'll come back to you personally"	CC of the same message
Paid for a charter	booking confirmation with the details	CC of the same message
Joined the crew list	welcome, with the unsubscribe line	CC of the same message
Bought a gift certificate	the certificate and how to redeem it	your own copy

You are CC'd on the guest's copy rather than sent a separate notification. That is deliberate, and it is a workaround. On 7 September a crew-list signup was accepted by Resend, returned an ID, and never arrived — the business's own sending domain silently stopped delivering while everything else kept working. Being on the guest's copy means you see the message they actually got, and it travels by the route that demonstrably works.

`bookings@thenautiyacht.com` **cannot receive mail**. The domain has no MX record, so anything sent to it is accepted and evaporates. It is still fine as a *from* address, and replies go to the Gmail address instead. This is unfixed — if anything ever assumes a reply can reach that address, it will fail.

How they sign off

Every one closes as **The Nauti Yacht LLC**, through a single setting, so the four cannot drift apart. A templated email is sent by the business, and a first name on it promises a named human is waiting at the other end of a reply.

The text messages are the opposite and stay that way: the review asks and the owed-charter messages read **"Austin & Brooke"**, because you paste those into a thread and send them from your own phone. The rule is who is really doing the thing, not which channel it went down.

Crew list

A phone number is now **required** to join. Corey signed up on 7 September with no number, which left email as the only way to reach him about a seat.

Joy also reads the signups on every run and puts new names in her status, because an email you cannot see fail is not a notification.

Why some things still say "Jarvis"

The Jarvis tab was retired on 3 September 2026. Six of its seven panels had become worse copies of what the Overview already showed; the seventh was the media pipeline, which moved to **Marketing -> Media Drafts** where it belonged. Speaking aloud moved to the top bar, where her last message appears with a speaker on it. Pearl is the one you talk to now. (There was briefly an **Enable Pearl** button; it is gone — see the top bar above.)

Nothing you can see says Jarvis any more. Four things under the floor still do, and they were left alone deliberately:

Still called	What it is	Why it stays
<code>JarvisTodo</code>	the database table behind The Board	renaming it means a migration against the live database to change a word
<code>JARVIS_SERVICE_KEY</code> / <code>x-jarvis-key</code>	how the crew scripts prove who they are	it is set in Vercel and in the secrets store; renaming means changing both in exact step or every agent stops being able to write
<code>/api/jarvis-todos</code>	the address The Board reads from	an address only has to be stable, not pretty
<code>Jarvis-Voice-UI</code>	the folder holding the crew scripts	hard-coded in fifteen scripts and nine agent briefs

Same rule as the task folders below: **rename what people read, leave what machines depend on**. If you see one of these in an error message, it is not a leftover anyone forgot.

Why the routines are named as they are

All eight now carry their crew member's name — `nauti-pearl`, `nauti-coral`, `nauti-siren`, `nauti-joy`, `nauti-penny`, `nauti-reef`, `nauti-shelly` and `nauti-nova`, alongside `crew-standup`. The renaming is finished; what follows is why it had to be done carefully, and it still applies to any future one.

A task ID cannot be renamed — only deleted and recreated — and connector approvals are stored **on the task**. Siren publishes through Blotato, so recreating her throws that approval away and her next run stops, waiting for a permission nobody is there to give. She is the one agent whose silence is invisible until an event has already gone unannounced.

So she was renamed only just after a publish run, when nothing was due and a **Run now** to bank the approval was a harmless no-op. Renaming her before one either loses that morning's posts to a prompt, or fires them at whatever hour you pressed the button.

Coral and Penny were renamed this way on 4 Sep 2026 — rename, **Run now**, approve once. Their approvals are Blotato and Gmail respectively. **If a routine ever starts skipping, a missing approval is the first thing to check.**

A folder name never tells you a cadence. `nauti-penny` runs daily and `nauti-shelly` runs weekly, and nothing in either name says so. **The schedule table above is what is true.**

Who checks whom

Everything routes through Pearl. Coral reports to Siren, and Siren reports to Pearl. The one loop that runs backwards is Coral auditing what Siren actually published the day before.

That loop exists because **Siren is the only agent that acts outside the business**. Everyone else proposes, so a mistake costs a line on the board; hers is live on three platforms. Joy, Reef and Nova are deliberately *not* checked by anyone — putting a reviewer in front of an agent that only proposes buys delay and no safety.

The Board (To-do List)

The centre panel on Overview. It is the shared workspace between you and the crew, not just your list — they write to it, read each other's items, and hand work over by naming another agent in an item.

Items are ranked **High / Medium / Low**, not listed by date. Crew items carry their own priority; the ones you typed yourself are ranked by what they say — money being held, security, legal and overdue all go high. Anything falling due within two days is promoted regardless of wording. Only High is expanded when the page loads; Medium and Low fold open when you want them.

Long items fold. The claim and its first two lines show; anything more sits behind + *N more lines*, and the dated notes behind *N notes*, the way they always have. One item was rendering as twenty bullets on 17 September and its own verdict — Coral writing "Both halves dealt with" — was at the bottom where nobody would see it. Nothing is dropped; it is a reading transform only.

Pearl keeps the board. She closes what the data shows is done, folds duplicates into the older entry, and rewrites anything Low that has sat untouched. She never edits or closes an item you wrote unless it is genuinely

finished.

Their status cards

Each agent files a status every morning **even on days she does not run**, so a card is never blank. A card that shows a date instead means she has not filed today.

Click her face to hear her read it. The small speaker mark on an avatar means she has something filed; clicking plays it in her own voice, clicking again stops her, and clicking again starts her from the beginning. Clicking a different agent stops whoever is talking first, so you never get two at once. Hovering enlarges the face, which is only there so you can see it properly.

Nothing in the console ever speaks on its own. If you did not click it, it will not make a sound.

Each click costs a little money. Speech is billed per character to ElevenLabs — roughly 350 characters a click, and \$10 buys about ten thousand. Clicking the same unchanged status twice is free; it replays from memory rather than being bought again.

A card reading **"Stopped mid-run"** means the run was killed partway, usually by a permission prompt nobody was there to answer. Open Routines and hit **Run now** to clear it.

When a card says an agent failed and she did not

A card is a row in the activity log, and that row carries one of four values. Three are what you would expect. The fourth is not a state at all:

Value	Means
running	started and has not closed yet
completed	finished
failed	genuinely went wrong
status	not a lifecycle state — the daily status card the standup files for each agent, even on days she does not run

`status` is why the table holds far more rows than there have been runs: eight agents filing one card a day. They are correct and should never be "cleaned up".

A card can report a failure that never happened. On 14 September Nauti Joy's card read failed for a day over a run that had succeeded four minutes earlier. The cause was in how agents record themselves: the script read its arguments by position, and the closing command passes an empty `""` for the title. PowerShell silently drops an empty argument to a program, so every later value slid one slot left and the status word landed in the title column.

Two things now stop it. The script refuses any status it does not recognise rather than writing a row it cannot vouch for, and closing a row uses named flags — `--status`, `--detail`, `--id` — so nothing has to survive being empty.

A card can also say an agent is STOPPED MID-RUN when she is not. On 17 September both Siren and Penny were reported stopped. Siren had published all three of that evening's posts ten minutes earlier with a URL recorded against every one. Both rows were status cards whose arguments had slid, so the word `status` sat in the task-title column and the status column kept its default `running` — the row never closes, and the console reports a death that did not happen. The logger now detects the slide, shifts every argument back and writes the row correctly, with a warning in the run transcript.

So if a card says failed, read its detail before believing it. A real failure explains what went wrong. An artefact usually has a status word sitting where the task title belongs, or says outright that it was a duplicate. Four such rows existed and were repaired on 16 September 2026; one of them had also had a genuine accounts-payable report overwritten on top of it, which was recovered from the duplicate before that duplicate was deleted.

The eight voices, and how each one writes

Every agent has her own voice and her own register, so you can tell who is speaking without looking at the name.

	Voice	How she writes
Pearl	Alice	Dry and unhurried. She tells you the thing you would rather not hear
Coral	Jessica	Keen, and hardest on her own work — she audits the queue she built
Siren	Charlotte	Flat and literal on purpose. She is the gate
Joy	Lily	Warm, and a little wounded on the guests' behalf
Penny	Matilda	Charming and teasing, openly proprietary about the books
Reef	Laura	The enthusiast. She is selling you an idea
Shelly	Sarah	Dry and sceptical — she asks why you pay for things you do not use
Nova	River	Sparse. She speaks rarely and it should feel like it cost her something

There is one hard limit on all of it: tone lives in the framing, never in the finding. Anything serious — money held for a charter that never ran, a booking at risk, anything with a legal, insurance or safety edge — is said straight, with no colour at all. That contrast is deliberate. A voice you enjoy, that goes flat the instant something is genuinely wrong, is a voice you look up for.

Their statuses are also now written to be **heard**, not only skimmed: full month names rather than "Sept", "4 September" rather than "9/4", "7pm" rather than "19:00", and no raw database ids or file paths. A checker runs whenever any of them files and warns her if she slips; it never blocks, because a badly written status is still worth more than no status.

What none of them may do

Every brief forbids the same things, and the shared protocol overrides any brief that disagrees:

- No agent writes to any table except the todo board and its own activity log.
- No agent contacts a guest. (The automatic DM reply is a Blotato platform feature, not a crew member — see *Messages answer themselves*.)
- No agent spends, refunds or changes a price.
- No agent publishes except Siren, and only drafts you have already scheduled.

The full rulebook lives at `C:\Users\immex\.claude\scheduled-tasks\_crew-protocol.md` . If a rule is not in there, it is not a rule.

Quick reference — where do I...

I want to...	Go to
See a new lead, from anywhere	Bookings
Log a Boatsetter/GetMyBoat charter	Bookings → Bookings
Record that a charter happened (<i>this logs the income too</i>)	Bookings → Bookings, set to completed
Block a day off	Bookings → Availability
Log a receipt	Money → Income & expenses
Check a charter's money is on the books	Money → Reconciliation
Pull tax numbers, or see per-charter and per-hour	Money → Tax Report
Add a photo to the public gallery	Marketing → Media
Approve or reschedule a social post	Marketing → Media Drafts
Say what is wrong with a post, or stop one	Marketing → Media Drafts → Discuss
See when a partly-booked day is actually taken	Bookings → Availability
Reply to a Facebook or Instagram comment	Marketing → Comments
Ask a past guest for a Google review	On the dock → Reviews (<i>phone</i>)
Change a price	Setup → Packages & pricing
Log engine hours after an outing	On the dock → Engine hours (<i>phone</i>)
Log fuel, or set service intervals	Boat → Maintenance

Things that are easy to get wrong

- **A tab with no number is not empty.** The number means "waiting on you".

- **Two income rows on one Boatsetter charter is correct** — boat leg and captain fee.
- **The review "Text it" link does nothing on a desktop.** Use a phone.
- **A completed booking with no price gets no income row.** Fill the price in.
- **Filenames lie about dates.** A file named `19:43` may have been *saved* then and shot at midday. Trust the timestamp inside the file, not the name.
- **The Lake Bryan charters of 6 and 13 June 2026 are under NDA.** No media from those may be posted, ever.
- **Party Cove / 2025-09-27 is footage from The Dam.** Mis-filed. The theme folder is the only thing that says where a clip was shot — the filenames are bare timestamps — so until it is moved, treat that date as The Dam. It had already become the opener and closer of a Party Cove compilation before you spotted it.
- **The blocked-media list is in `write-tags.js`, not only in the tags file.** That script REBUILDS the list wholesale rather than adding to it, so a block recorded only in `_media-tags.json` disappears the next time anybody runs the tagger — silently, with searches simply starting to return the file again. On 18 September the script held three entries and the live list held seven. Four files would have been unblocked, two of them with a confirmed topless woman.
- **A block on a file is not always a block on every second of it.** The 20 September 2025 glow video is blocked for a figure at 4.5–6.5s; an excerpt from 42–54s was reviewed and cleared. Where that has happened the `doNotUse` entry names the seconds and says who cleared what, so the block keeps meaning something instead of being re-argued every time.
- **One glow clip is withdrawn for nudity:** `Boatz and Glowz / 2025-09-20 First run - promo event / 20250920_211850_303ad61a.mp4`. The whole clip, not just the part that was used, because it cannot be reviewed reliably at any size that makes screening practical.
- **"Don't post" does not delete.** It moves the draft to Rejected, and *Back to review* brings it back. Deleting is a separate button behind a confirm.
- **Six hashtags is a failed post, not a style note.** The publisher rejects it outright and the draft still looks scheduled. Five, on every platform.
- **A photo on Instagram or TikTok is untested.** Both accounts have only ever published video. Facebook takes either.
- **"I posted it myself" is only for posts you published by hand.** Siren marks her own and records the link.
- `bookings@thenautiyacht.com` **cannot receive mail.** Fine to send *from*, never somewhere to expect a reply.
- **A reply under a DM may not be yours.** Most are answered automatically within seconds. Do not read one as a conversation you have had.
- **Boatz & Glowz ropes off at 5, not 7, and runs to around midnight.** Check-in 4:30, seven hours. It moved twice on 17 September — the finish first, then the departure — and every time it moves, fourteen captions, two DM automations, the package row, the public FAQ and the live bookings all have to follow. **The times live in `lib/glowEvent.js`:** `GLOW_CHECK_IN_TIME`, `GLOW_START_TIME`, `GLOW_RETURN_TIME`, and `GLOW_TIMING_LINE`, which is the whole sentence already written. Change them there, never in the copy.
- **The glow page counts its own seats**, and says the number twice — beside the Reserve button and again on the Seats card. Both read the same count, so they cannot disagree. "31 seats" and "seats left" are worked out from the bookings on the night, never typed in — every seat somebody is expecting counts as taken,

including unpaid inquiries and the crew riding free, because a public page that promises a seat twice ends at the ramp. It refreshes every 60 seconds.

- **A post that has already published cannot be edited by anything here.** Blotato has no route for it — four posts had to be corrected by hand in each app when the time moved. Facebook shows "Edited"; Instagram does not; TikTok cannot be edited at all.
- **The console cannot see a reply you sent from the Facebook app.** Press *I answered this elsewhere* or the thread keeps its red flag forever.
- **A red flag on a message thread means a machine answered something it should not have.** Read what went out before you reply.
- **The bank balance is typed in, not fetched.** If it says 9d ago in amber, nobody has looked in nine days and the figure is a guess with a date on it.
- **A Wells Fargo warning does not stop you filing the row.** File it if it is true. The thing to act on is whatever is still charging that account.